

FEDERATED LEARNING

Presented by Vudit Khandelwal

Overview

- Introduction
- Flower
- Implementation
- Flower
- Results
- Results
- Conclusion
- Flower

Introduction

Centralized Learning is a paradigm of Machine Learning that has been widely explored and utilized. But, due to the massive amount of computations required, it requires High Performance Computers and Cloud Computers for its training procedure, along with the need to collect huge amounts of data from devices across the globe.

Federated Learning provides an interesting approach to this problem of machine learning. Rather than collecting the data and bringing it to server, it performs training on many edge devices and aggregates the changes in model parameters to train the model.

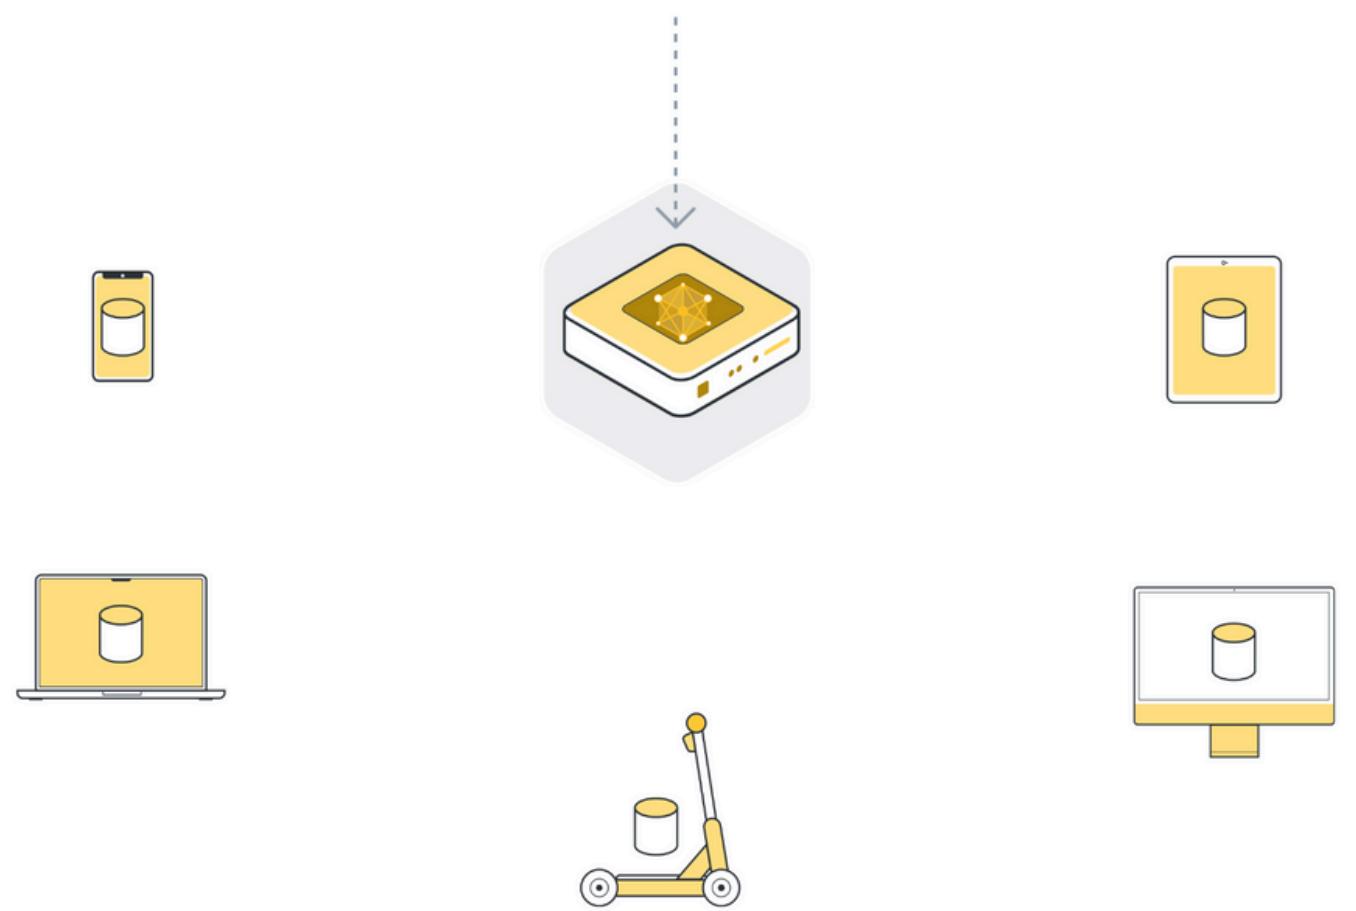
Problem

- We have one server, and N clients. Each client is an edge device that will have a locally trained model. Only these changes in parameters of model are transmitted back to the server.

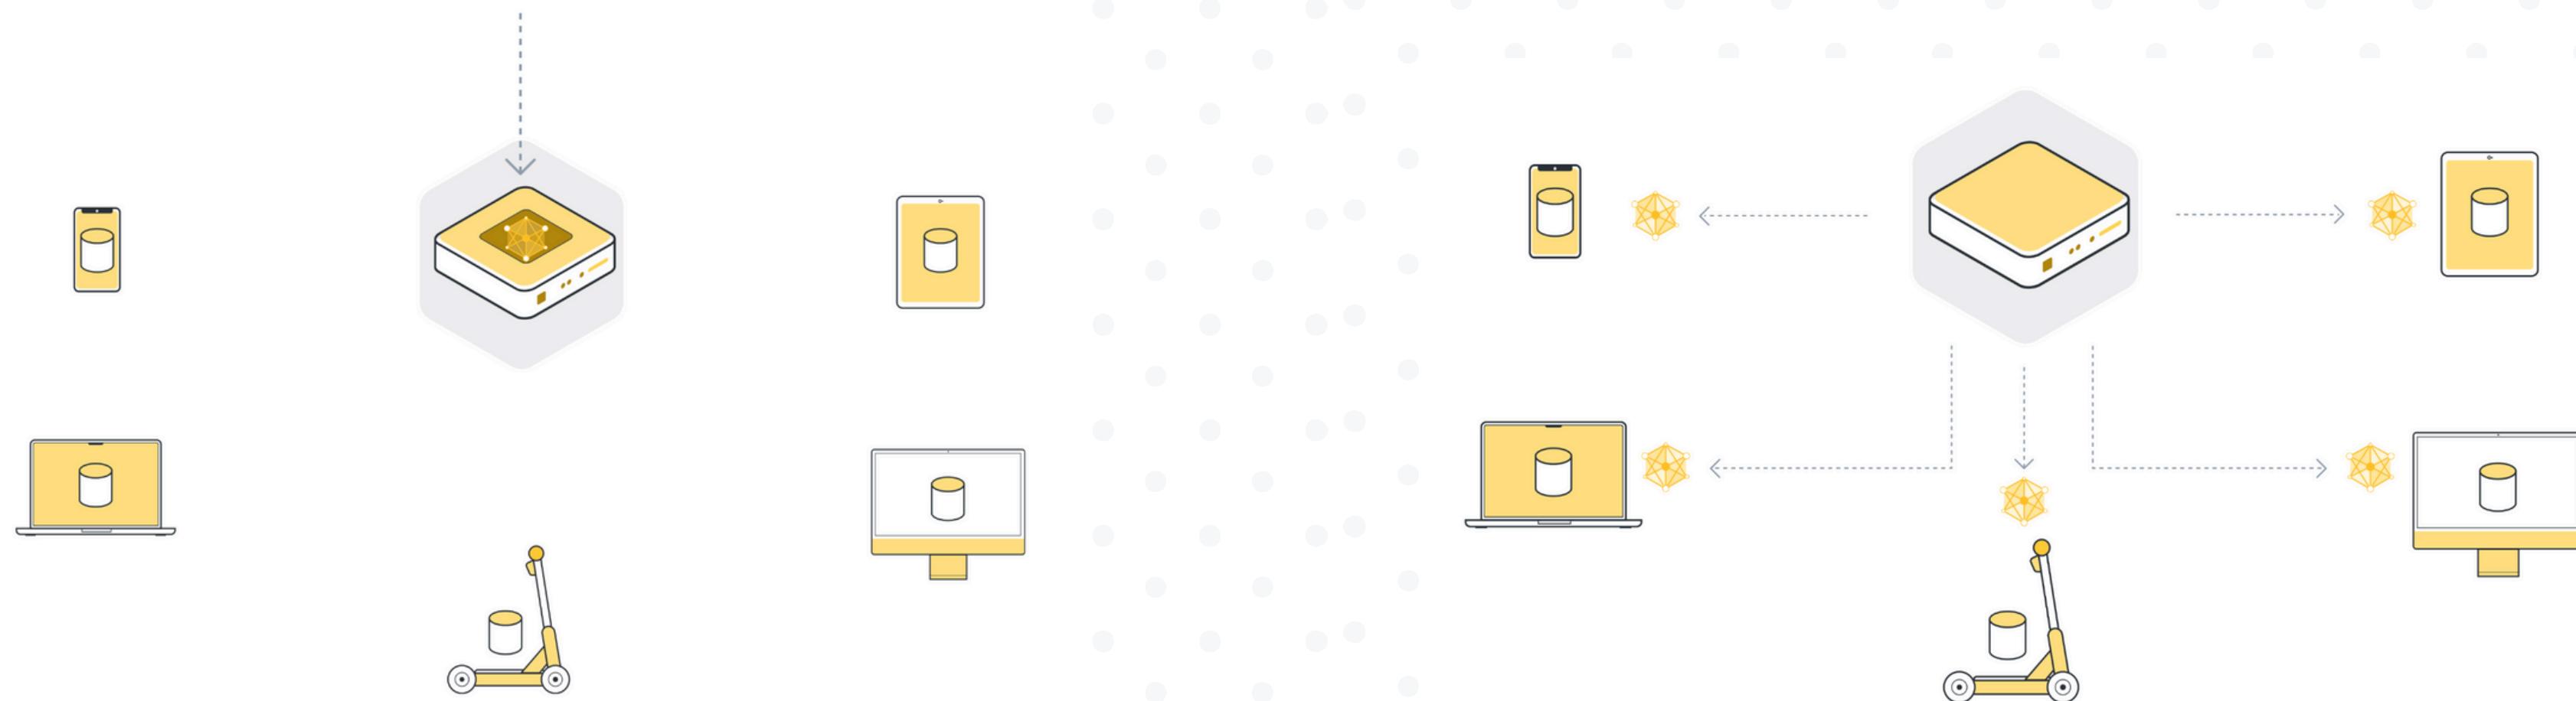
Thus, federated learning has advantages of:

- Privacy, data stays at edge devices
- Diverse data from multiple devices
- Data security,

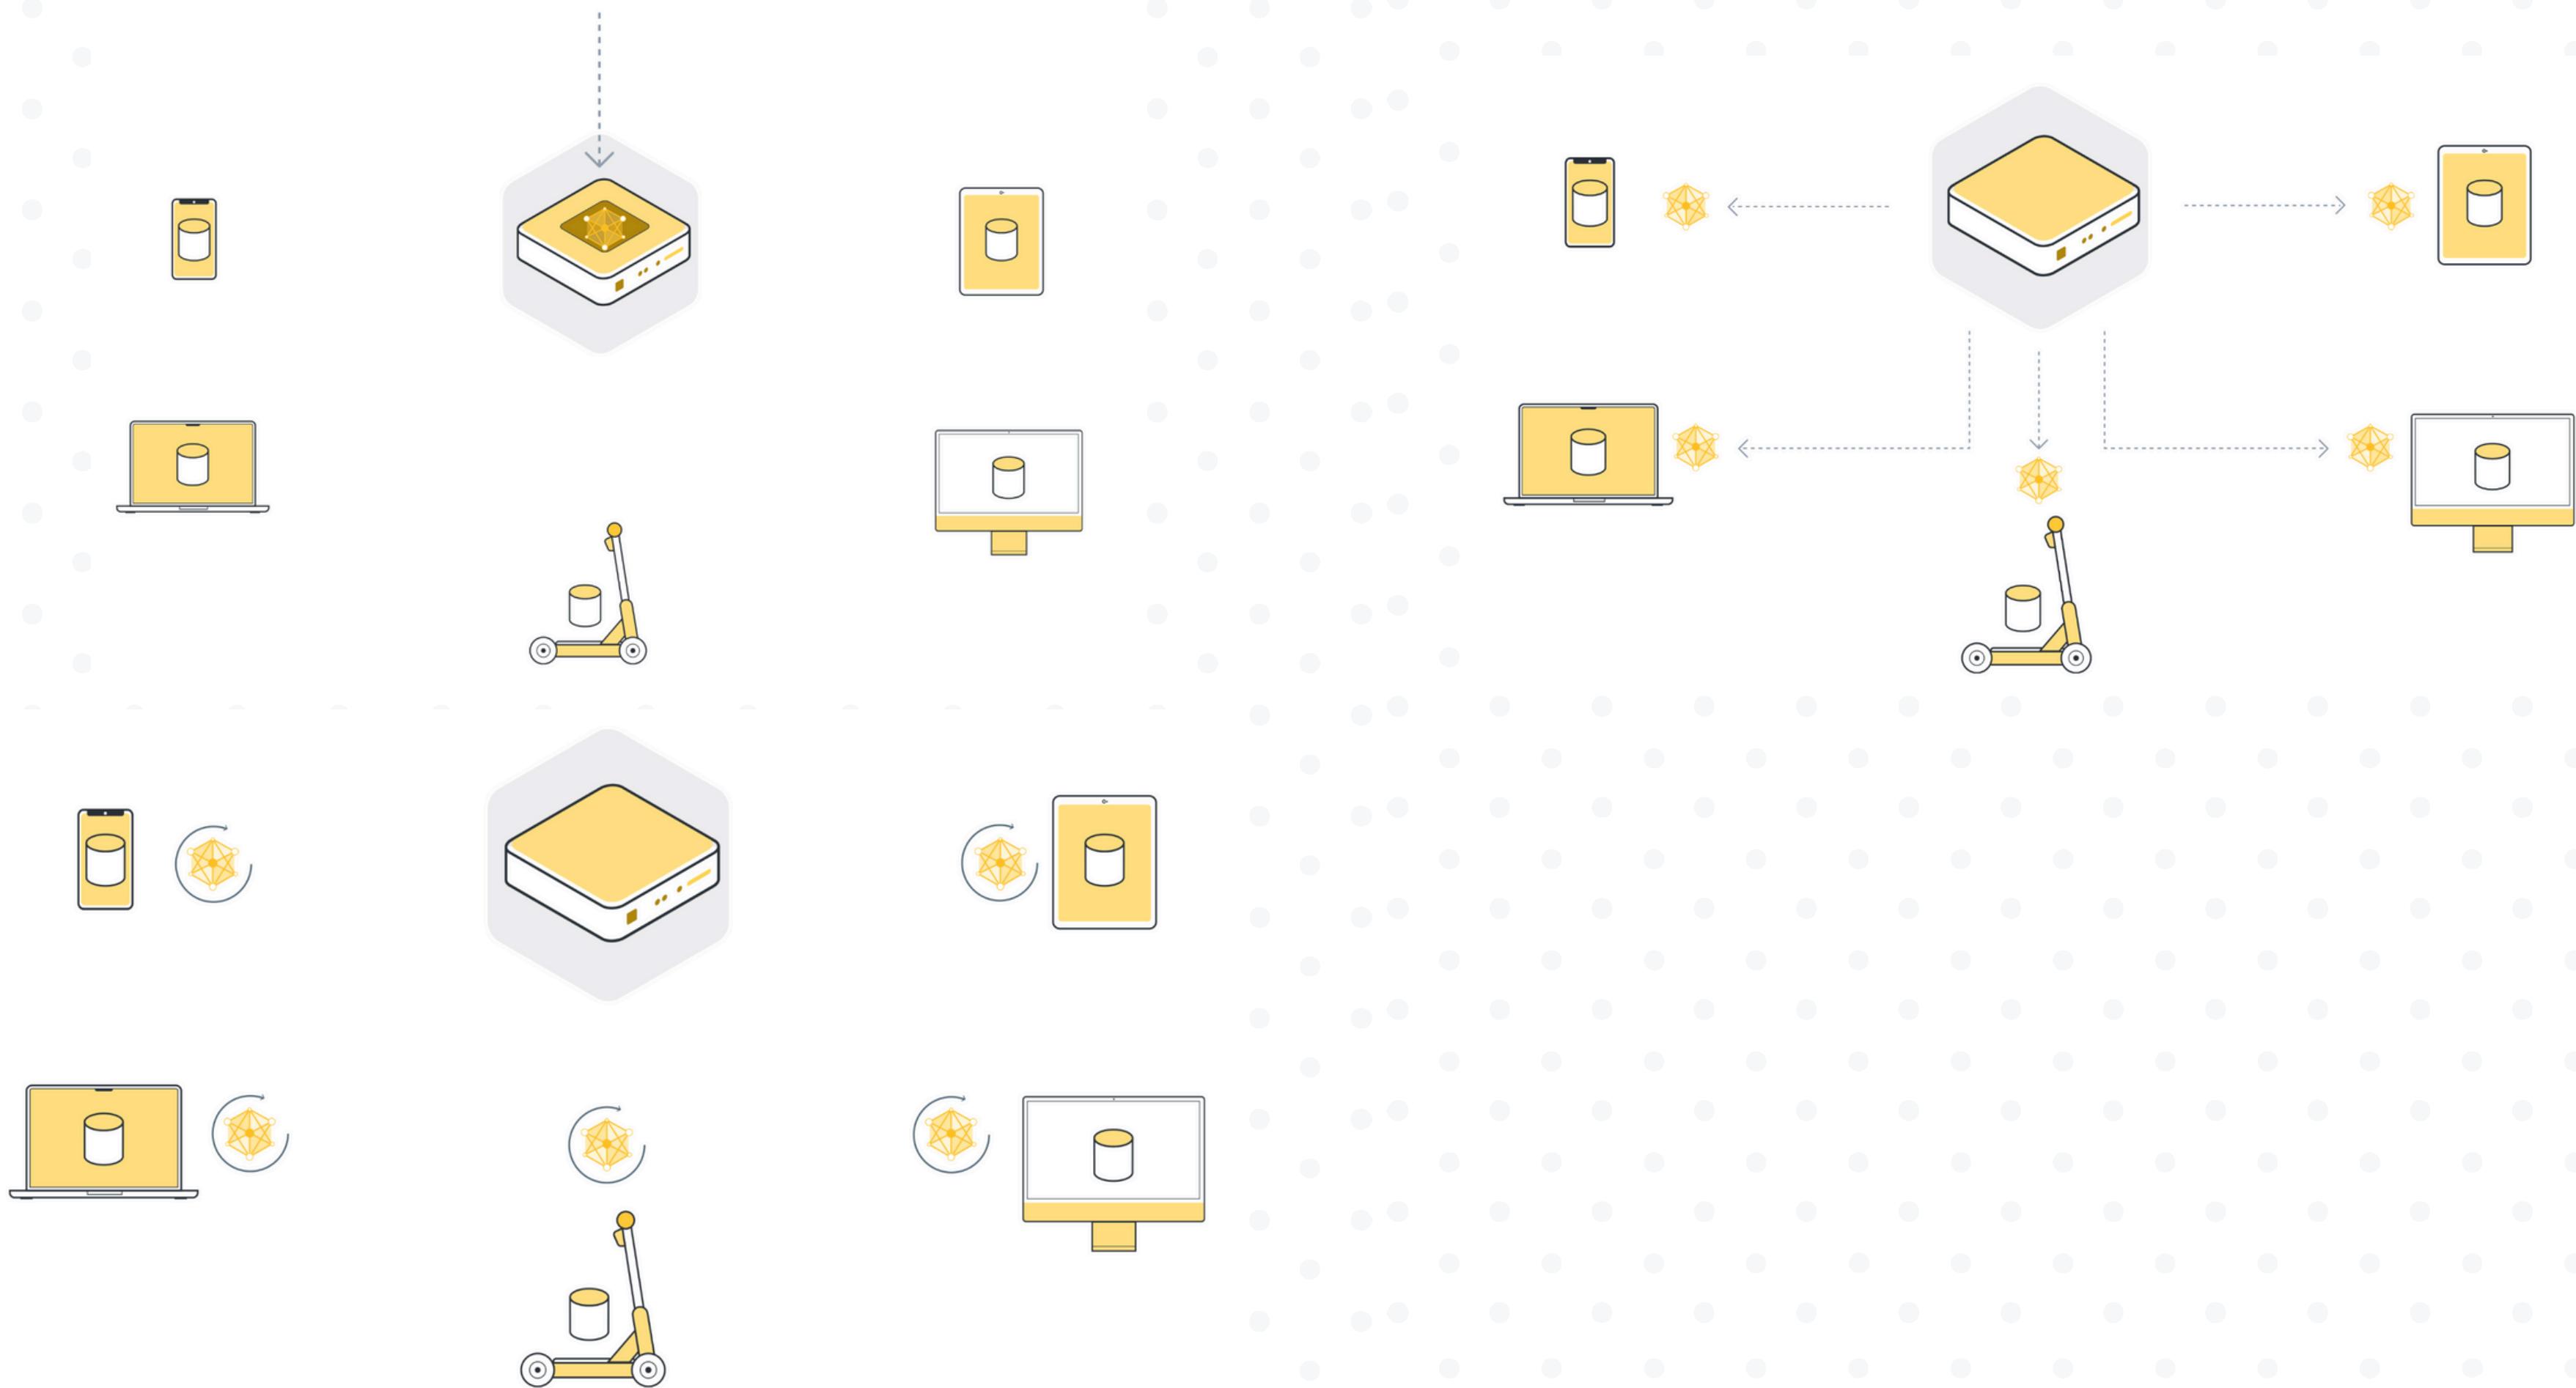
Step 1. Initialize Global model



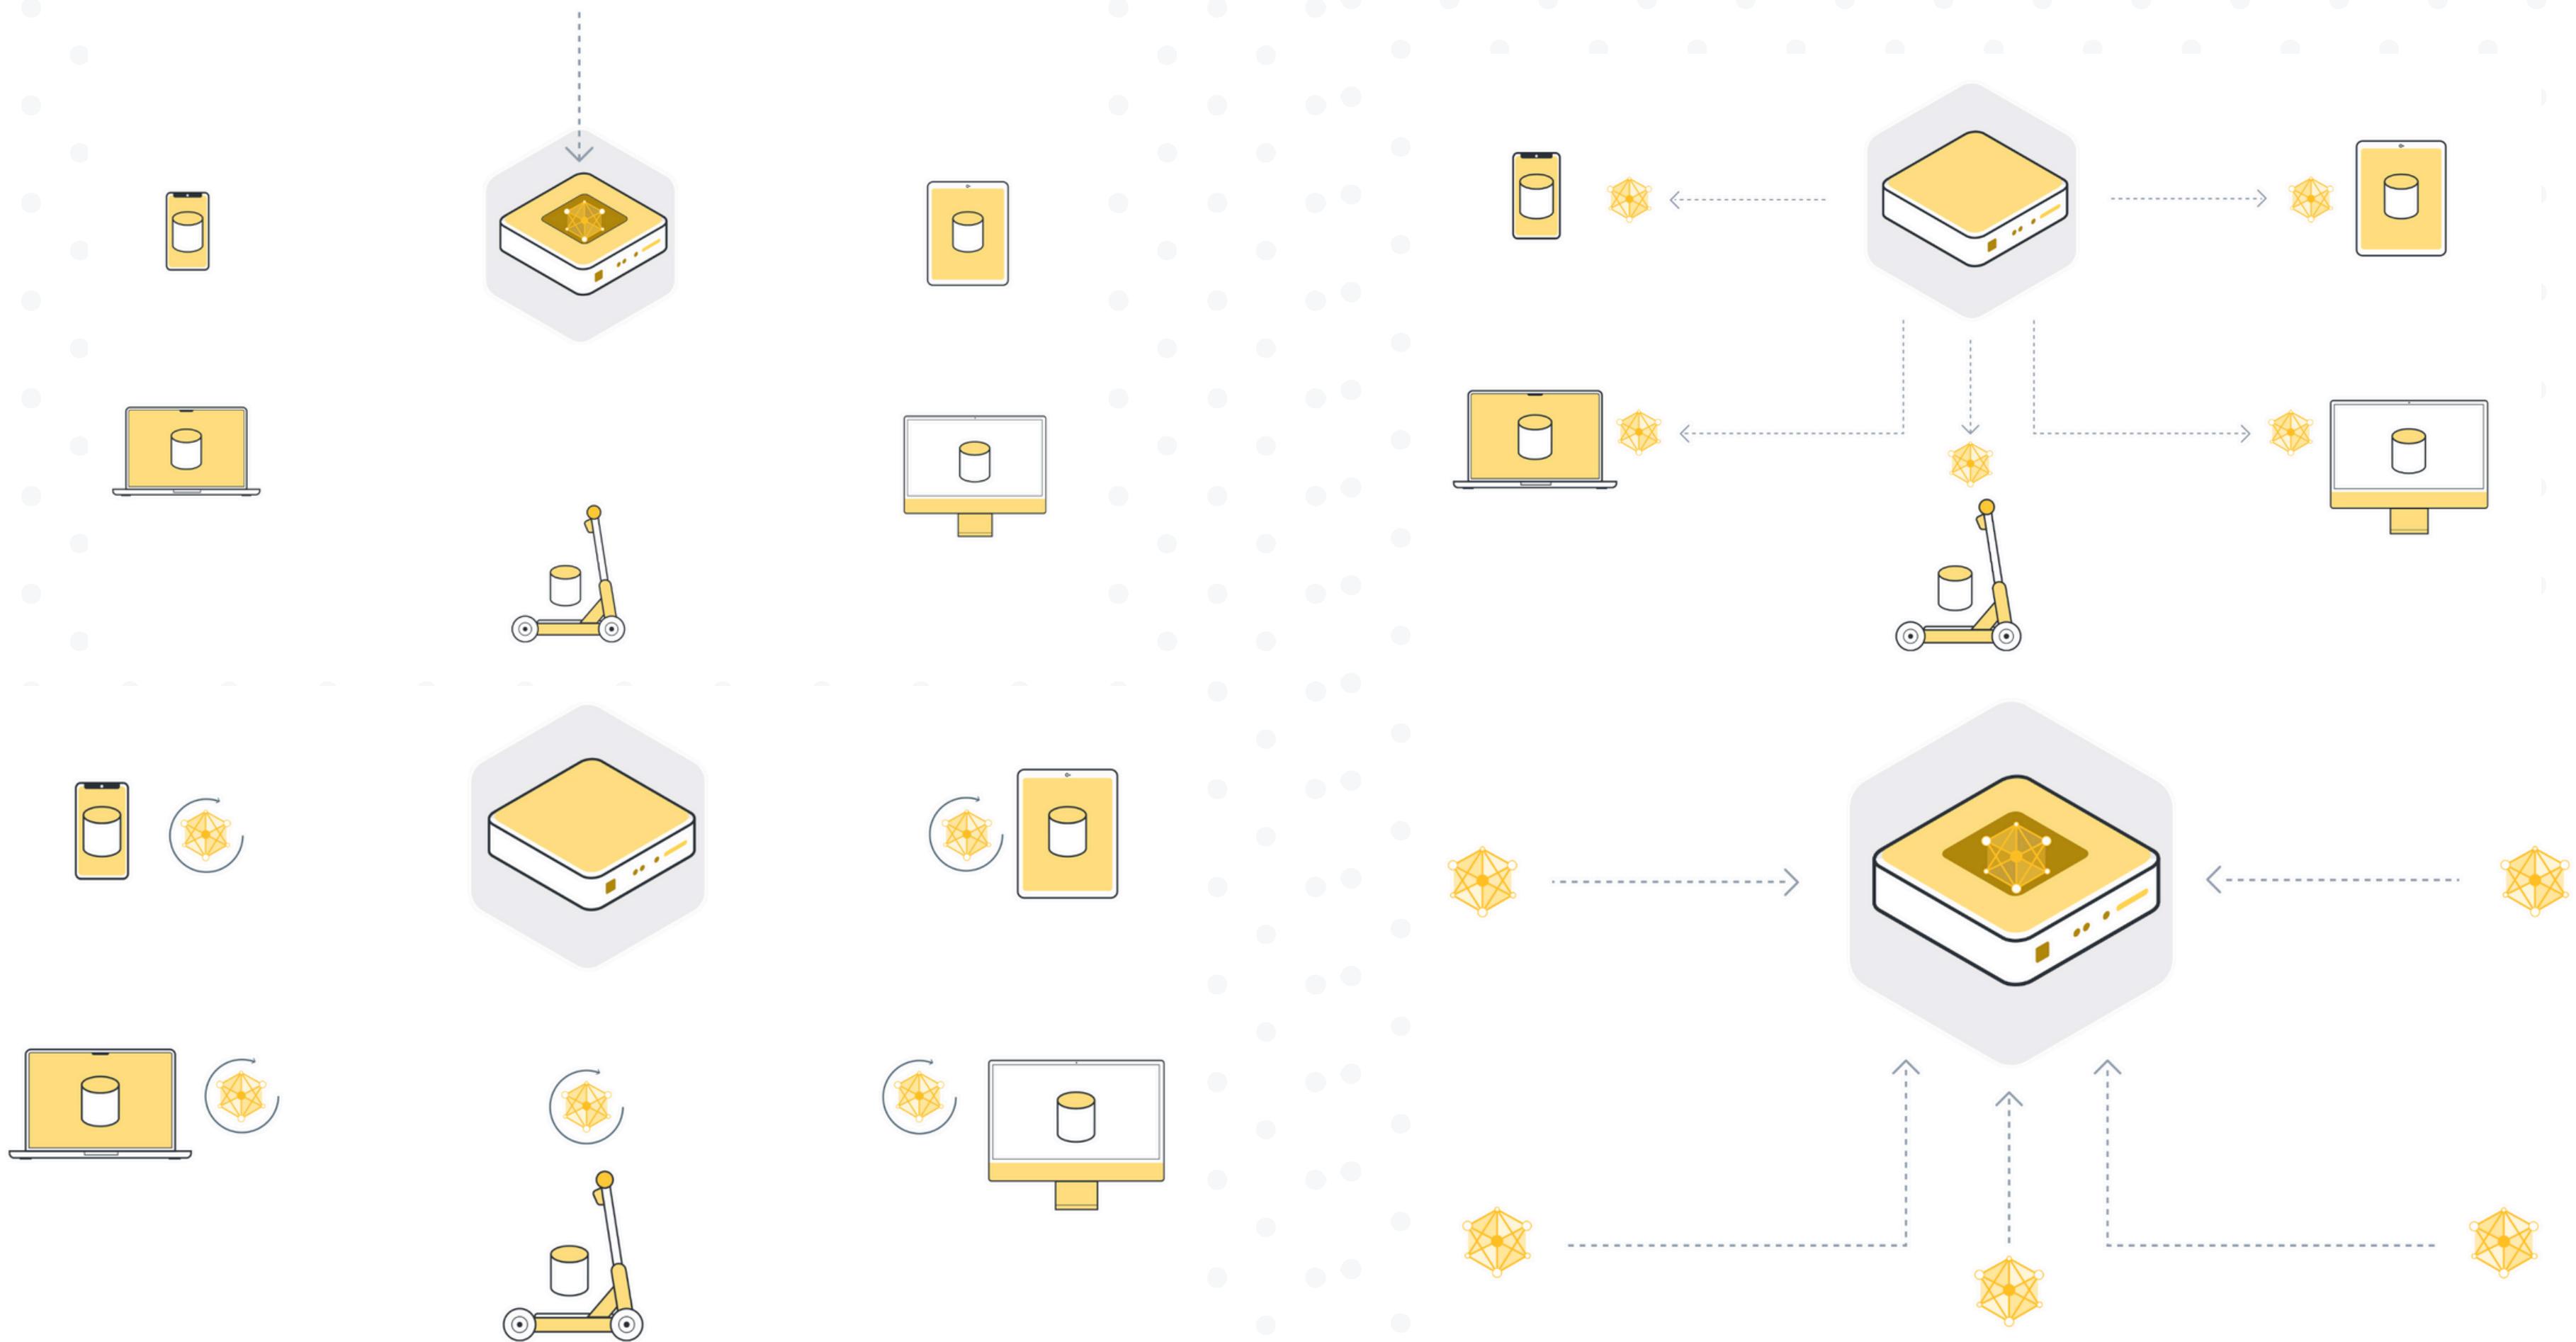
Step 2. Send model to a number of clients



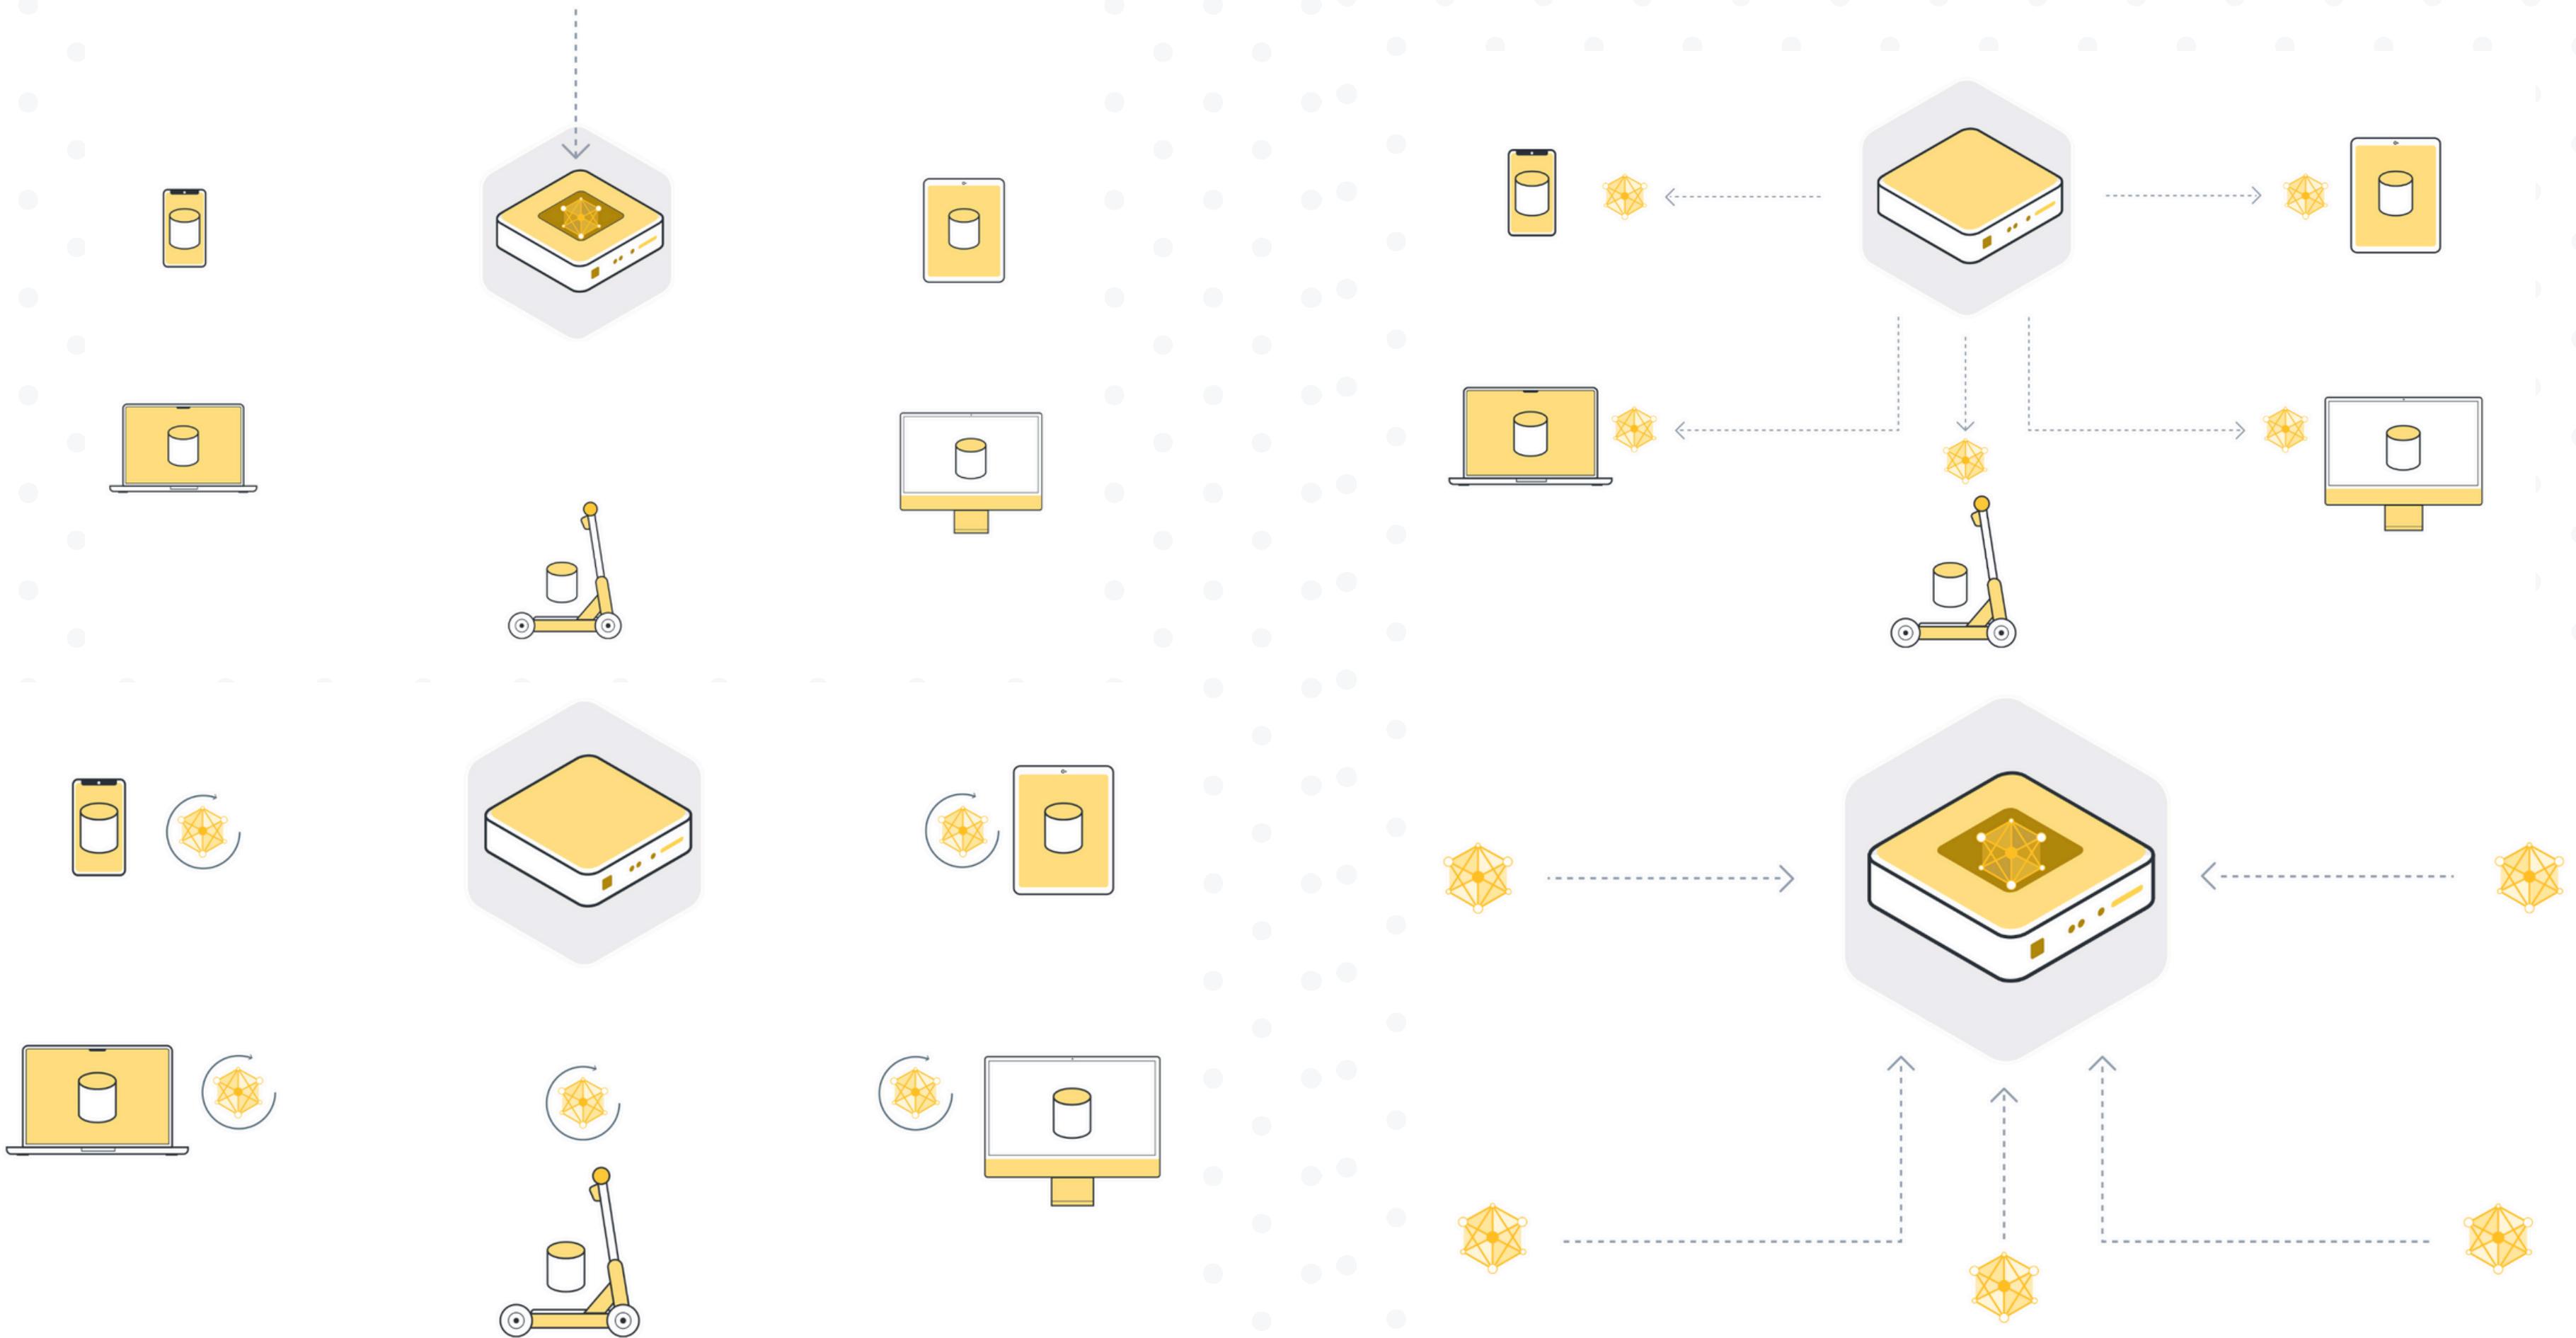
Step 3. Train the model locally on edge devices



Step 4. Return model updates to the server



Step 5. Repeat steps 1-4 for N iterations



Flower: A Friendly Federated AI Framework



[Website](#) | [Blog](#) | [Docs](#) | [Conference](#) | [Slack](#)

license [Apache-2.0](#) PRs [welcome](#)  Framework [passing](#) downloads [578k](#) Docker Hub [flwr](#) Chat [Slack](#)

Flower (`flwr`) is a framework for building federated AI systems. The design of Flower is based on a few guiding principles:

- **Customizable:** Federated learning systems vary wildly from one use case to another. Flower allows for a wide range of different configurations depending on the needs of each individual use case.
- **Extendable:** Flower originated from a research project at the University of Oxford, so it was built with AI research in mind. Many components can be extended and overridden to build new state-of-the-art systems.
- **Framework-agnostic:** Different machine learning frameworks have different strengths. Flower can be used with any machine learning framework, for example, [PyTorch](#), [TensorFlow](#), [Hugging Face Transformers](#), [PyTorch Lightning](#), [scikit-learn](#), [JAX](#), [TFLite](#), [MONAI](#), [fastai](#), [MLX](#), [XGBoost](#), [Pandas](#) for federated analytics, or even raw [NumPy](#) for users who enjoy computing gradients by hand.
- **Understandable:** Flower is written with maintainability in mind. The community is encouraged to both read and contribute to the codebase.

Objective

1. Using the Flower framework, simulate the process of Federated Learning.
2. Change the model and see changes
3. Change the dataset and show that federated learning still works
4. Compare the performance of federated with centralized learning
5. Finally work on prototype portraying FL

1.Implementing FL using Flower

A simple CNN model
that can be trained to for
classification of images

```
def get_model(width: int = 32, height: int = 32, num_channels: int = 3) -> Any:  
  
    model = models.Sequential([  
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),  
        layers.BatchNormalization(),  
        layers.Conv2D(32, (3, 3), activation='relu'),  
        layers.MaxPooling2D((2, 2)),  
        layers.Dropout(0.25),  
  
        layers.Conv2D(64, (3, 3), activation='relu'),  
        layers.BatchNormalization(),  
        layers.Conv2D(64, (3, 3), activation='relu'),  
        layers.MaxPooling2D((2, 2)),  
        layers.Dropout(0.25),  
  
        layers.Flatten(),  
        layers.Dense(128, activation='relu'),  
        layers.Dropout(0.5),  
        layers.Dense(10, activation='softmax')  
    ])  
    model.compile("adam", "sparse_categorical_crossentropy", metrics=["accuracy"])  
    return model
```

CIFAR-10: an image classification dataset

```
def load_data(
    partition_id: int, num_partitions: int
) -> tuple[np.ndarray, np.ndarray, np.ndarray, np.ndarray]:
    """Load data with Flower Datasets (CIFAR-10)."""
    # Only initialize `FederatedDataset` once
    global fds
    if fds is None:
        partitioner = IidPartitioner(num_partitions=num_partitions)
        fds = FederatedDataset([
            dataset="uoft-cs/cifar10",
            partitioners={"train": partitioner},
        ])
    partition = fds.load_partition(partition_id, "train")
    partition.set_format("numpy")

    # Divide data on each node: 80% train, 20% test
    partition = partition.train_test_split(test_size=0.2, seed=42)
    x_train, y_train = partition["train"]["img"] / 255.0, partition["train"]["label"]
    x_test, y_test = partition["test"]["img"] / 255.0, partition["test"]["label"]

    return x_train, y_train, x_test, y_test
```

The **SERVER** function defines the aggregation fn used for FL, here used FedAvg

```
def server_fn(context: Context) -> ServerAppComponents:
    """Construct components that set the ServerApp behaviour."""

    # Read from config
    num_rounds = context.run_config["num-server-rounds"]

    model = get_model()
    ndarrays = get_parameters(model)
    global_model_init = ndarrays_to_parameters(ndarrays)

    # Define strategy and the custom aggregation function for the evaluation metrics
    strategy = FedAvg(
        evaluate_metrics_aggregation_fn=average_metrics,
        initial_parameters=global_model_init,
    )
    config = ServerConfig(num_rounds=num_rounds)

    return ServerAppComponents(strategy=strategy, config=config)

# Create ServerApp
app = ServerApp(server_fn=server_fn)
```

The **CLIENT** function performs the training(fitting) of the model

```
# Define Flower client
class FlowerClient(NumPyClient):
    # pylint: disable=too-many-arguments
    def __init__(self, model, x_train, y_train, x_test, y_test):
        self.model = model
        self.x_train = x_train
        self.y_train = y_train
        self.x_test = x_test
        self.y_test = y_test

    def fit(self, parameters, config):
        self.model.set_weights(parameters)
        self.model.fit(
            self.x_train, self.y_train, epochs=1, batch_size=32, verbose=False
        )
        return self.model.get_weights(), len(self.x_train), {}

    def evaluate(self, parameters, config):
        self.model.set_weights(parameters)
        loss, accuracy = self.model.evaluate(self.x_test, self.y_test, verbose=False)
        y_pred = self.model.predict(self.x_test, verbose=False)
        y_pred = np.argmax(y_pred, axis=1).reshape(
            -1, 1
        ) # MobileNetV2 outputs 10 possible classes, argmax returns just the most probable

        acc, rec, prec, f1 = eval_learning(self.y_test, y_pred)
        output_dict = {
            "accuracy": accuracy, # accuracy from tensorflow model.evaluate
            "acc": acc,
            "rec": rec,
            "prec": prec,
            "f1": f1,
        }
        return loss, len(self.x_test), output_dict
```

Experiments & Results

Dataset: CIFAR10

Model: CNN

20 rounds of FL, 10 Clients

Accuracy: 62%

f1 score:0.619

```
viditkh@pop-os: ~/Desktop/Federated Learning/testing/custom-metrics
INFO : round 20: 1.1264501571655274
INFO : History (metrics, distributed, evaluate):
INFO : {'acc': [(1, 0.3383000000000005),
INFO : (2, 0.4133),
INFO : (3, 0.4468),
INFO : (4, 0.4705000000000003),
INFO : (5, 0.4974),
INFO : (6, 0.5084),
INFO : (7, 0.5239),
INFO : (8, 0.5355000000000001),
INFO : (9, 0.5450999999999999),
INFO : (10, 0.5541),
INFO : (11, 0.5628),
INFO : (12, 0.5742999999999999),
INFO : (13, 0.5783999999999999),
INFO : (14, 0.5885),
INFO : (15, 0.5926),
INFO : (16, 0.5974),
INFO : (17, 0.6068),
INFO : (18, 0.6112),
INFO : (19, 0.6162),
INFO : (20, 0.619)],
INFO : 'accuracy': [(1, 0.338299986410141),
INFO : (2, 0.4133000162124636),
INFO : (3, 0.446799994754791),
INFO : (4, 0.470499983310699),
INFO : (5, 0.497400006914139),
INFO : (6, 0.508400021219254),
INFO : (7, 0.523900022411346),
INFO : (8, 0.535500019073486),
INFO : (9, 0.545099975204468),
INFO : (10, 0.554100008583069),
INFO : (11, 0.56279996137619),
INFO : (12, 0.574300030517578),
INFO : (13, 0.578400039100647),
INFO : (14, 0.588499990463257),
INFO : (15, 0.592599939441681),
INFO : (16, 0.597399977111817),
INFO : (17, 0.606799958992004),
INFO : (18, 0.611199928951264),
```

2. Changing the image classification model used

Dataset: CIFAR10
Model: MobileNetV2
5 rounds of FL, 10 Clients

Accuracy: 30.1%
f1-score:0.3

```
viditkh@pop-os: ~/Desktop/Federated Learning/testing/custom-metrics
[INFO : [SUMMARY]
[INFO : Run finished 5 round(s) in 184.08s
[INFO : History (loss, distributed):
[INFO :     round 1: 2.1119752645492555
[INFO :     round 2: 2.015006446838379
[INFO :     round 3: 1.9740293979644776
[INFO :     round 4: 1.9505069851875305
[INFO :     round 5: 1.9351823449134826
[INFO : History (metrics, distributed, evaluate):
[INFO : {'acc': [(1, 0.2614999999999995),
[INFO :             (2, 0.2794999999999997),
[INFO :             (3, 0.2890000000000003),
[INFO :             (4, 0.2957),
[INFO :             (5, 0.3013)],
[INFO : 'accuracy': [(1, 0.261500009536743),
[INFO :                 (2, 0.279499986886978),
[INFO :                 (3, 0.2889999856948855),
[INFO :                 (4, 0.2957000171661376),
[INFO :                 (5, 0.3012999816417696)],
[INFO : 'f1': [(1, 0.2614999999999995),
[INFO :             (2, 0.2794999999999997),
[INFO :             (3, 0.2890000000000003),
[INFO :             (4, 0.2957),
[INFO :             (5, 0.3013)],
[INFO : 'prec': [(1, 0.2614999999999995),
[INFO :                 (2, 0.2794999999999997),
[INFO :                 (3, 0.2890000000000003),
[INFO :                 (4, 0.2957),
[INFO :                 (5, 0.3013)],
[INFO : 'rec': [(1, 0.2614999999999995),
[INFO :                 (2, 0.2794999999999997),
[INFO :                 (3, 0.2890000000000003),
[INFO :                 (4, 0.2957),
[INFO :                 (5, 0.3013)]}
```

3. Changing the dataset used

Dataset: FashionMNIST
Model: CNN
10 rounds of FL, 50 Clients

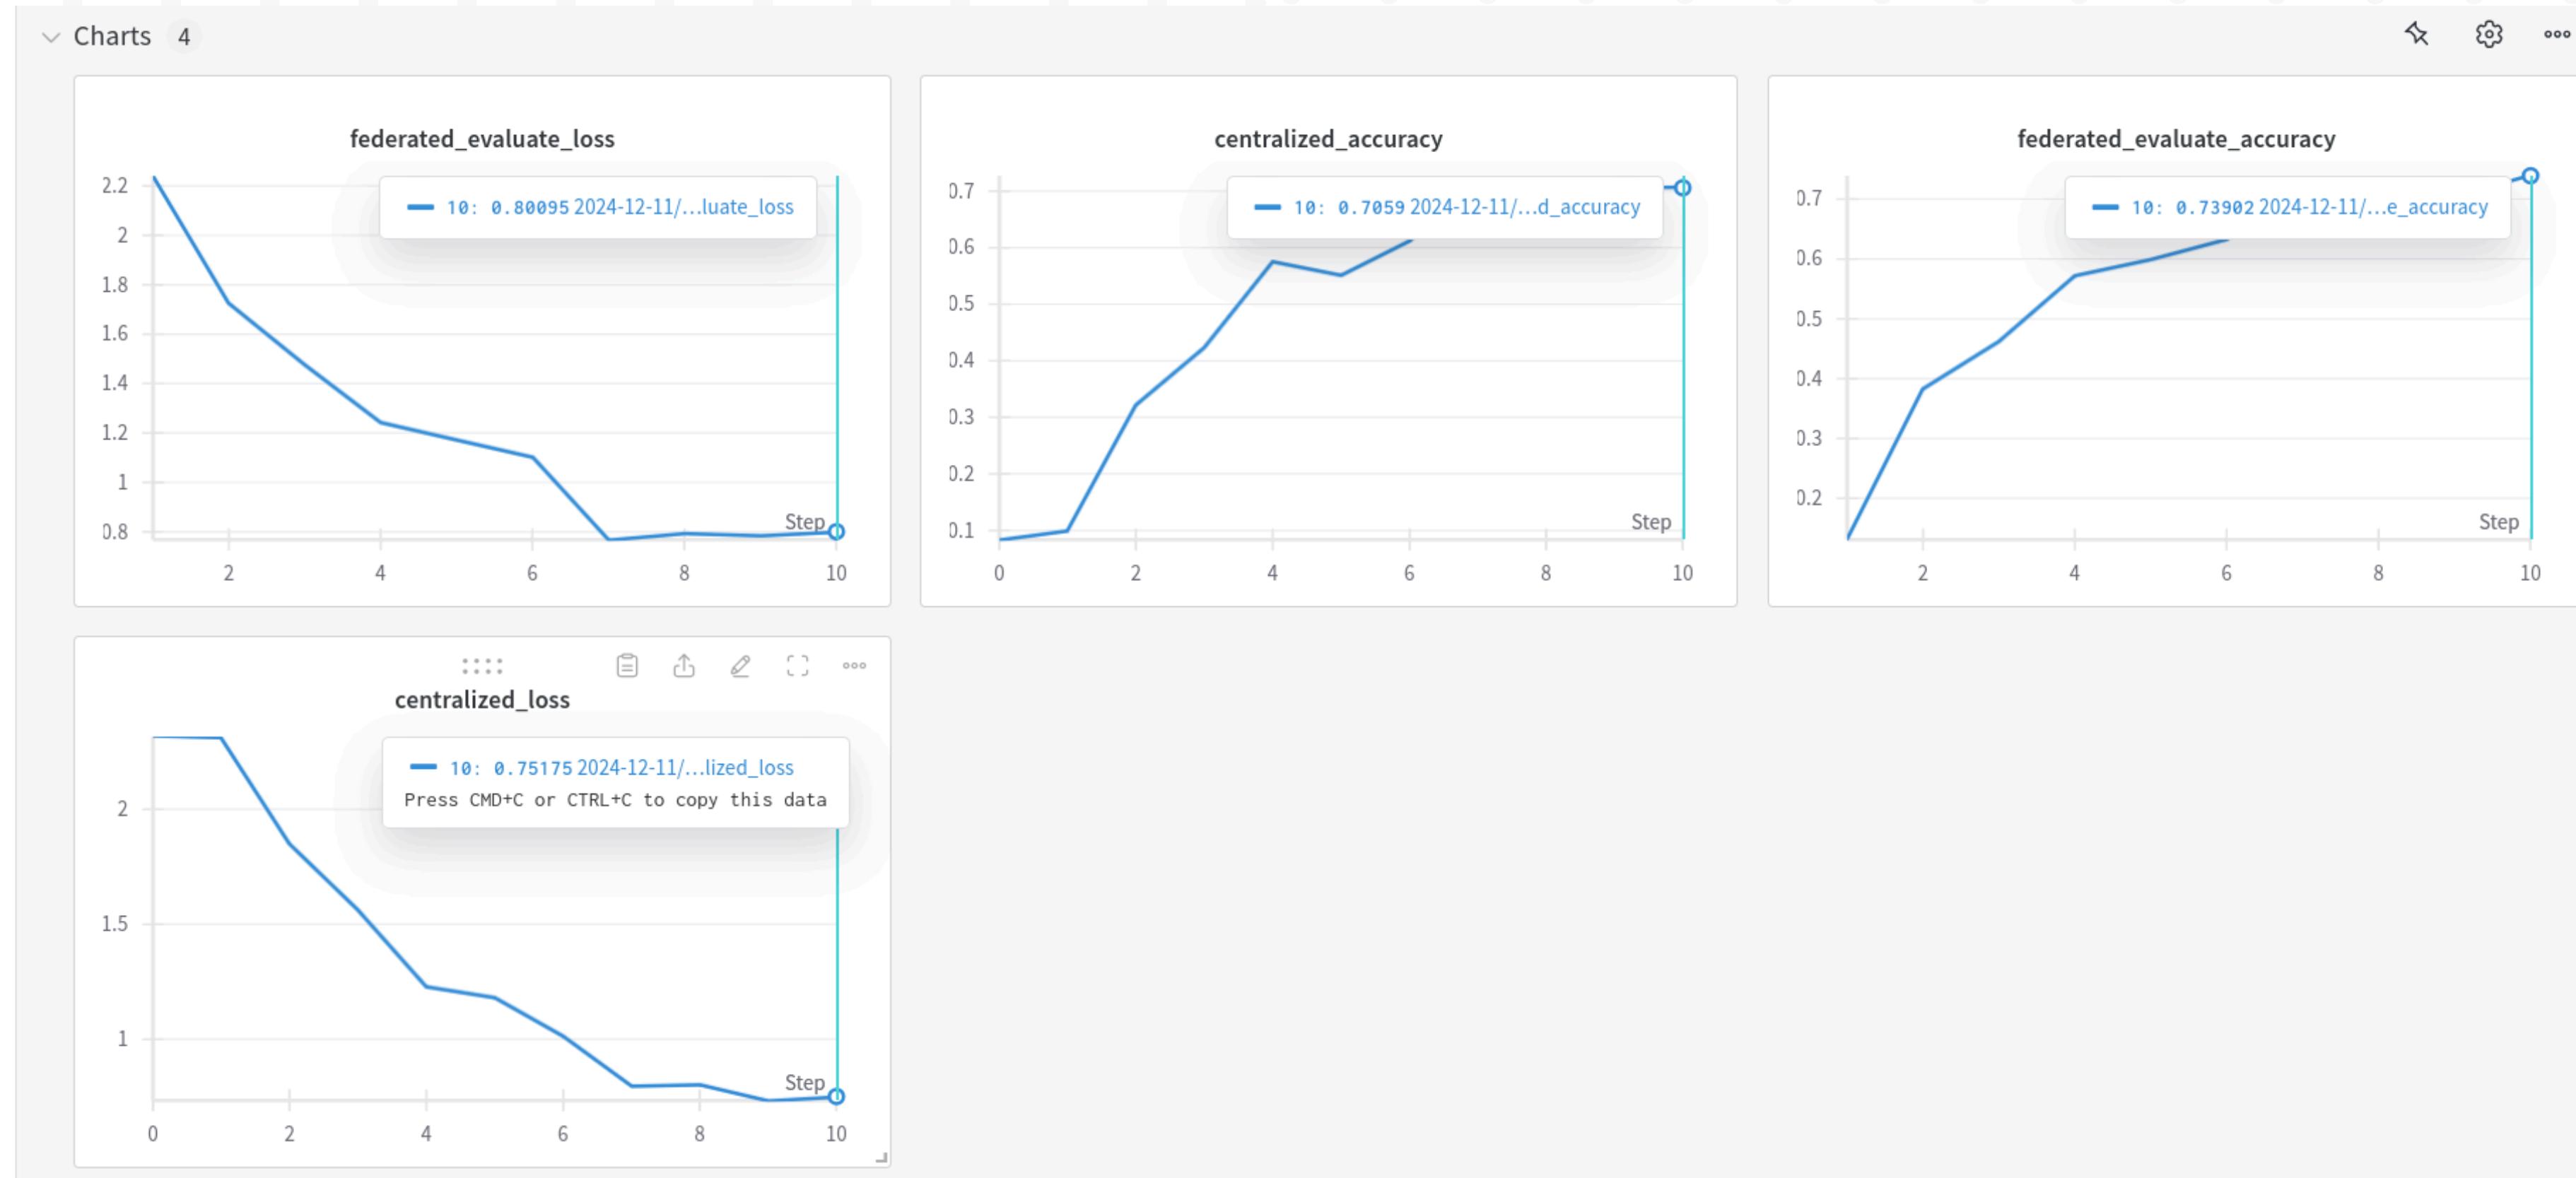
CL Accuracy: 70%
FL Accuracy: 73%

```
Workspaces Applications
F
viditkh@pop-os: ~/Desktop/FederatedLearning

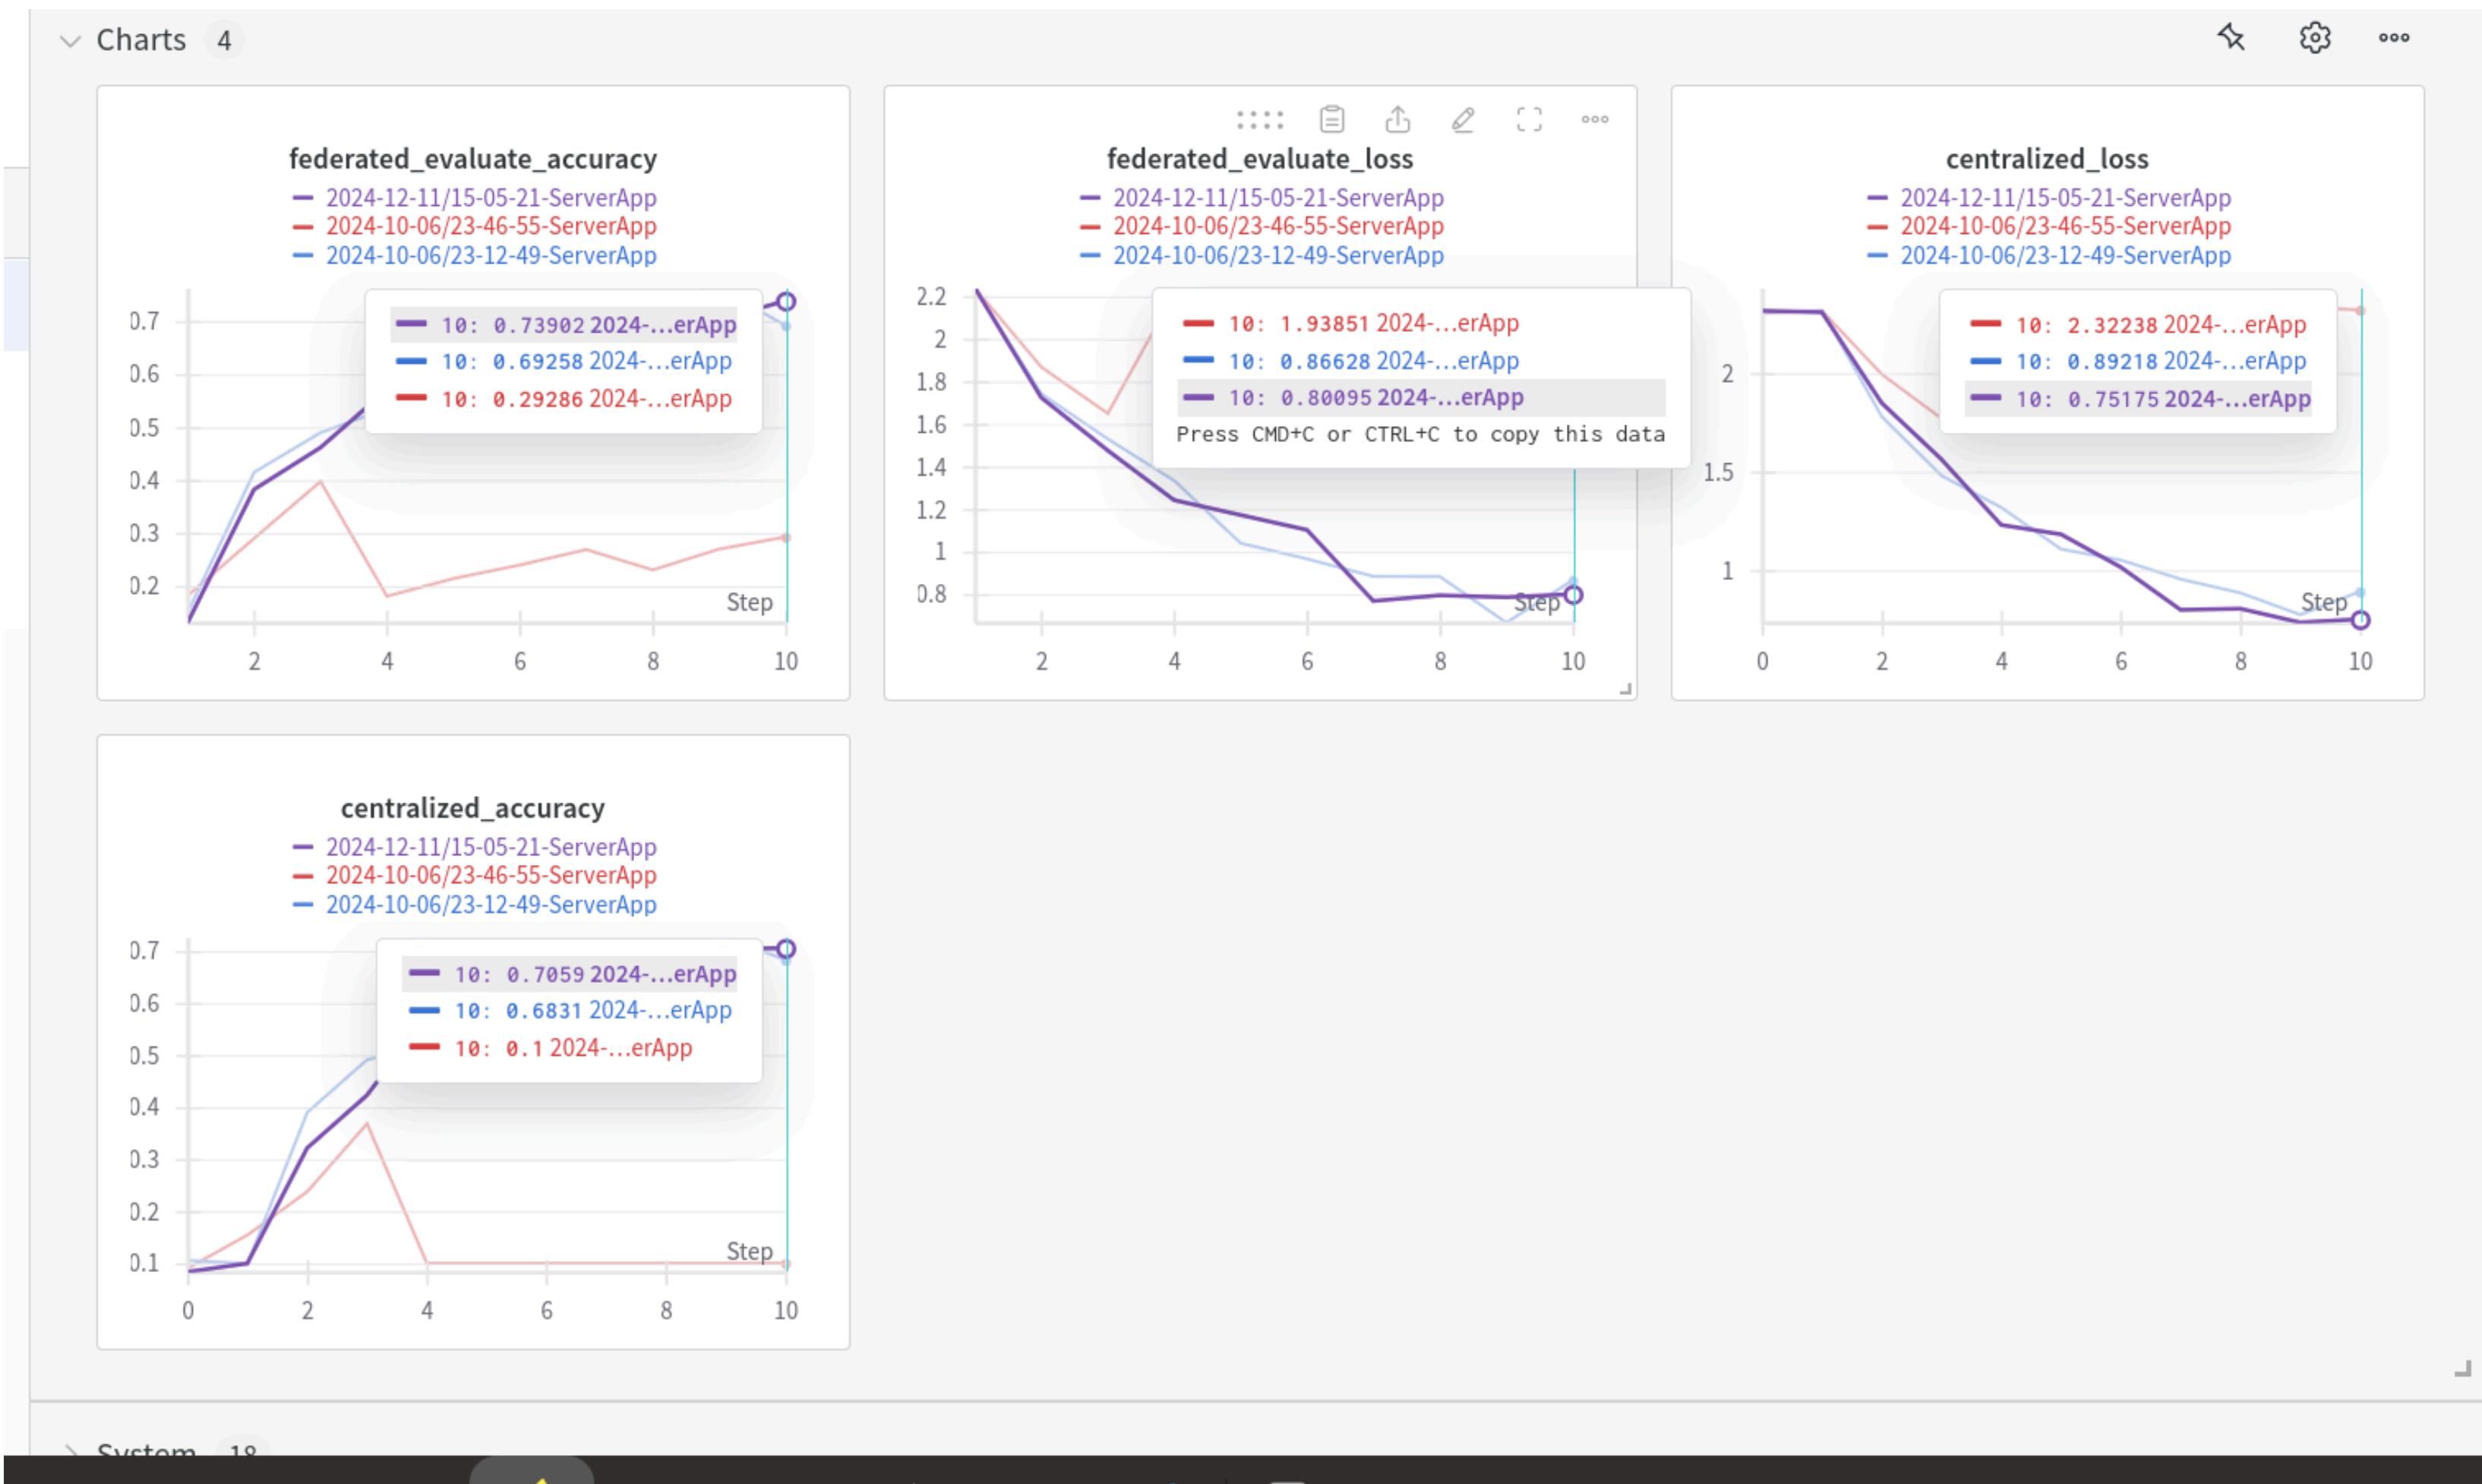
INFO : round 6: 1.1029793154703555
INFO : round 7: 0.7686879378185325
INFO : round 8: 0.7951544313273512
INFO : round 9: 0.7862665864435313
INFO : round 10: 0.8009479276932358
INFO : History (loss, centralized):
INFO : round 0: 2.3163422372775337
INFO : round 1: 2.3090347862853027
INFO : round 2: 1.8487442305293709
INFO : round 3: 1.5627494524843015
INFO : round 4: 1.2293441794550837
INFO : round 5: 1.1820796179695252
INFO : round 6: 1.015322785027111
INFO : round 7: 0.7989300596066558
INFO : round 8: 0.8047124817729377
INFO : round 9: 0.7347460435792661
INFO : round 10: 0.7517502451666628
INFO : History (metrics, distributed, evaluate):
INFO : {'federated_evaluate_accuracy': [(1, 0.13111221614453838),
INFO : (2, 0.38279472007282656),
INFO : (3, 0.4617934403349616),
INFO : (4, 0.5713610586011342),
INFO : (5, 0.5983683715092564),
INFO : (6, 0.6318974724101104),
INFO : (7, 0.7334256534533495),
INFO : (8, 0.7253837804310745),
INFO : (9, 0.7037466843501327),
INFO : (10, 0.7390176409546869)]}
INFO : History (metrics, centralized):
INFO : {'centralized_accuracy': [(0, 0.0838),
INFO : (1, 0.1),
INFO : (2, 0.3217),
INFO : (3, 0.423),
INFO : (4, 0.5745),
INFO : (5, 0.5505),
INFO : (6, 0.6111),
INFO : (7, 0.7272),
INFO : (8, 0.6974),
INFO : (9, 0.7045),
INFO : (10, 0.7059)]}
```

4. Comparing performance of FL v/s CL

Graphs Plotted accuracy v/s Rounds



Graphs comparing FL v/s CL



Conclusion

- Using the Flower Framework for Federated Learning Simulation, implemented server and client simulating FL.
 - Then achieved the above results in different scenarios like
 - Changing the dataset
 - Changing the ML Model used
 - Finally, compared performance of FL with Centralized Learning
-
- Now next step is to work on developing prototype for the simulations performed. Have learnt ways to connect devices wirelessly, will continue to work on this.

Larana University | 2025

Thank You

Presented by Aaron Loeb

Logistic Reg for MNIST Dataset

Accuracy: 83.9%

```
Workspaces Applications
viditkh@pop-os: ~/Desktop/Federated Learning/sklearn-logreg-mnist

INFO : [ROUND 1]
INFO : configure_fit: strategy sampled 5 clients (out of 10)
INFO : aggregate_fit: received 5 results and 0 failures
WARNING : No fit_metrics_aggregation_fn provided
INFO : fit progress: (1, 1.3041110766418718, {'accuracy': 0.6769}, 8.102675664999879)
INFO : configure_evaluate: strategy sampled 10 clients (out of 10)
INFO : aggregate_evaluate: received 10 results and 0 failures
WARNING : No evaluate_metrics_aggregation_fn provided
INFO :
INFO : [ROUND 2]
INFO : configure_fit: strategy sampled 5 clients (out of 10)
INFO : aggregate_fit: received 5 results and 0 failures
INFO : fit progress: (2, 0.5813381218367087, {'accuracy': 0.8197}, 17.217699679999896)
INFO : configure_evaluate: strategy sampled 10 clients (out of 10)
INFO : aggregate_evaluate: received 10 results and 0 failures
INFO :
INFO : [ROUND 3]
INFO : configure_fit: strategy sampled 5 clients (out of 10)
INFO : aggregate_fit: received 5 results and 0 failures
INFO : fit progress: (3, 0.5290841542400078, {'accuracy': 0.8397}, 20.32259933499995)
INFO : configure_evaluate: strategy sampled 10 clients (out of 10)
INFO : aggregate_evaluate: received 10 results and 0 failures
INFO :
INFO : [SUMMARY]
INFO : Run finished 3 round(s) in 22.34s
INFO : History (loss, distributed):
INFO :     round 1: 1.2923177606135794
INFO :     round 2: 0.6141979962266947
INFO :     round 3: 0.5612966216039441
INFO : History (loss, centralized):
INFO :     round 0: 2.3025850929940455
INFO :     round 1: 1.3041110766418718
INFO :     round 2: 0.5813381218367087
INFO :     round 3: 0.5290841542400078
INFO : History (metrics, centralized):
INFO :     {'accuracy': [(0, 0.098), (1, 0.6769), (2, 0.8197), (3, 0.8397)]}
INFO :
```